

# **Bitter Melons**

## A Review-Aggregator Database

### Design and Implementation Report

**Manish Tawade**  
Student ID: 24205067  
University College Dublin

# Contents

Contents . . . . .	i
List of Figures . . . . .	ii
List of Tables . . . . .	iii
1 Introduction . . . . .	1
2 Goals & Objectives . . . . .	3
3 Design & ER Diagrams . . . . .	5
3.1 Conceptual Overview . . . . .	5
3.2 Design Motivations . . . . .	6
4 Normalization . . . . .	8
4.1 First and Second Normal Form . . . . .	8
4.2 Third Normal Form . . . . .	8
4.3 Boyce–Codd Normal Form . . . . .	9
5 Implementation Details . . . . .	10
5.1 Schema Definition . . . . .	10
5.2 Procedural Elements . . . . .	11
5.3 Views . . . . .	13
5.4 Sample Data Population . . . . .	17
6 Example Queries . . . . .	18
6.1 Q1 – HoneyDon’t to HoneyDew Turnaround . . . . .	18
6.2 Q2 – Critic Scale Breakdown . . . . .	19
6.3 Q3 – Critics Who Specialize in One Genre . . . . .	20
6.4 Q4 – Top Actor per Franchise by Average Melon Score . . . . .	22
7 Conclusions . . . . .	23
7.1 Key Findings . . . . .	23
7.2 Strengths & Weaknesses . . . . .	24
7.3 Future Work . . . . .	24

## List of Figures

1	Entity-Relationship diagram with new tables. . . . .	6
2	Test 1 – score update lowers <code>sweetness_pct</code> . . . . .	11
3	Test 2 – deleting a low-profile critic cascades and decrements outlet totals. . . . .	12
4	Test 3 – duplicate review rejected by <code>UNIQUE</code> key. . . . .	12
5	Test 4 – brand-new review auto-generates cache rows. . . . .	12
6	Top 10 movies released in the last 24 months that earned a HoneyDew certification. . . . .	13
7	The five critics with the highest sweetness percentages (minimum 25 reviews). . . . .	14
8	All outlets ranked by their average sweetness percentage. . . . .	14
9	Sean Connery’s filmography ordered by feature sweetness percentage. . . . .	15
10	The ten action-genre films with the highest sweetness percentages. . . . .	15
11	The five genres with the highest average sweetness percentages. . . . .	16
12	Franchises with $\geq 3$ films ranked by average sweetness, with a side-by-side comparison of the Marvel Cinematic Universe and James Bond. . . . .	16
13	Q1 result – features rebounding from HoneyDon’t to HoneyDew. . . . .	19
14	Q2 result – critic review counts and pass rates by scale. . . . .	20
15	Q3 result – critics specializing in a dominant genre. . . . .	21
16	Q4 result – top-rated actor per franchise. . . . .	22

## List of Tables

1	Foreign-key cascade rules. . . . .	5
2	Normal-form compliance by table. . . . .	8
3	Stored program objects. . . . .	11
4	Summary of reporting views . . . . .	13
5	Design appraisal. . . . .	24

# 1 Introduction

In today’s streaming-first world, audiences grapple with an overwhelming flood of new releases. Whether it’s indie darlings on boutique platforms or summer blockbusters in multiplexes, viewers crave a reliable signal: “Is this worth my time?” Review-aggregator portals like Rotten Tomatoes® and Metacritic [4] have cemented their place in marketing campaigns and social media chatter by collapsing dozens of critic reviews into a single, easy-to-digest score. With just a *Fresh/Rotten* badge or a percentage score, these services steer billions of ticket-buying and streaming decisions each year.

Yet for many smaller departments and research groups, building a bespoke aggregator remains out of reach. The Department’s legacy `movies` database already captures titles, genres, principal cast, and franchise metadata—but lacks any facility for ingesting or analyzing professional reviews. As a result, instructors and students rely on manual spreadsheets or external APIs, impeding hands-on exploration of advanced SQL concepts like normalization and trigger-driven pipelines.

## User Pain-Points & Market Gaps.

- *Heterogeneous Rating Scales:* Five-star systems, percentage scores, ten-point scales, and thumbs-up/thumbs-down all coexist. Converting them into a coherent signal is non-trivial.
- *Performance at Scale:* Computing live aggregates over millions of rows can introduce unacceptable latency in a teaching environment.
- *Data Integrity & Versioning:* Ensuring each critic’s single vote per title—and preserving historical snapshots—requires careful constraint design.

**Project Vision.** This capstone effort, code-named **Bitter Melons**, transforms the legacy schema into a full-stack review-aggregator. By modeling publications (`outlets`) and reviewers (`critics`), normalizing diverse rating schemes, and enforcing strict referential integrity, we deliver a platform that:

- Normalizes a variety of scoring systems into a binary UP/DOWN recommendation.
- Guarantees one review per critic per feature via UNIQUE constraints and careful FK rules.
- Powers sub-second dashboard queries through three precomputed cache tables (`title_stats`, `critic_stats`, `outlet_stats`), refreshed by row-level triggers.
- Ships seven turnkey SQL views—leaderboards, rankings, genre heat-maps, franchise analyses—so that analysts spend time on insights, not JOIN syntax.
- Provides 1 300+ synthetic reviews (2008–2024) over 100+ features for reproducible

testing and trend exploration.

**Technical Challenges & Innovations.** To hit both pedagogical and performance targets, Bitter Melons employs:

- *Surrogate Keys & Controlled Denormalization:* Balancing normalization rigor with query efficiency.
- *Trigger-Driven Workflows:* Encapsulating cache refresh logic in a single `sp_refresh_stats` procedure.
- *Temporal Data Handling:* Archiving previous verdicts to support “what-if” analyses and time-series views.

Taken together, these design decisions yield a robust, future-proof schema that’s ideal for both real-world analytics and classroom exploration. In Section 2, we lay out our explicit goals and measurable objectives; Section 5.1 then constructs the conceptual and logical ER models, which we drive through each normal form up to BCNF [2]. Subsequent sections detail views, procedural elements, and example queries that showcase Bitter Melons’ full capabilities.

## 2 Goals & Objectives

This section translates our high-level vision for the Bitter Melons platform into concrete, measurable goals. Each objective drives a core piece of functionality in the review-aggregator architecture, and will be used later to validate our schema design, implementation, and performance.

### Goals:

- **Critic Metadata Integration.** *Objective:* Model every review event by recording the publication (`outlets`) and reviewer (`critics`), timestamped and tagged as “top” or “regular.” *Success criteria:*
  - All 40 synthetic critics and 10 outlets load without errors.
  - Queries such as `SELECT * FROM critics WHERE status='top'` return the expected subset in under 50 ms.
- **Multi-Scheme Rating Support.** *Objective:* Normalize disparate scoring systems (five-star, percentage, ten-point, thumbs-up/thumbs-down) into a unified binary UP/DOWN verdict via `rating_scales` and the `fn_to_recommendation()` function. *Success criteria:*
  - At least four rating schemes map correctly to UP/DOWN for all edge values (e.g. 2.5/5 stars = DOWN, 60% = UP).
  - Calling `fn_to_recommendation()` on any raw score executes in under 1 ms.
- **Data Integrity & Uniqueness.** *Objective:* Enforce a strict one-review-per-critic-per-feature rule using a UNIQUE constraint, and maintain referential integrity with carefully chosen CASCADE and RESTRICT actions on foreign keys. *Success criteria:*
  - Attempts to insert duplicate reviews are rejected by the DBMS.
  - Deleting an outlet with active reviews is blocked (RESTRICT), while removing a critic automatically cleans up orphaned cache entries (CASCADE).
- **High-Performance Dashboards.** *Objective:* Precompute key aggregates—total reviews and positivity rates—into three cache tables (`title_stats`, `critic_stats`, `outlet_stats`), refreshed by a single stored procedure `sp_refresh_stats()` invoked via row-level triggers. *Success criteria:*
  - Dashboard queries against any `*_stats` table return within 100 ms, even after bulk inserts of 1 300 reviews.
  - Trigger overhead adds no more than 5% latency to single-row insert/delete

operations.

- **Analyst-Ready Views.** *Objective:* Deliver seven out-of-the-box SQL views—covering release leaderboards, critic/outlet rankings, genre sentiment heat-maps, franchise performance, and cast insights—so end users can query with a simple `SELECT *` call. *Success criteria:*
  - Each view executes in under 200 ms.
  - Sample use-cases (e.g. “top 5 most positive genres in 2023”) run correctly without additional JOIN logic.
- **Reproducible Sample Data.** *Objective:* Seed the database with 4 rating schemes, 10 outlets, 40 critics, and roughly 1 300 reviews spanning 2008–2024, using a bulk-insert script that accepts a “sweetness bias” parameter for consistent replay. *Success criteria:*
  - Running the seed script twice with the same bias yields identical data distributions.
  - The overall positivity rate can be tuned between 30%–70% via a single script argument.
- **Normalization to BCNF.** *Objective:* Drive every table through 1NF, 2NF, 3NF, and finally BCNF—documenting each decomposition step and justifying any controlled denormalizations for performance. *Success criteria:*
  - A normalization report shows zero partial or transitive dependencies after each NF stage.
  - All BCNF violations (if any) are explicitly documented with rationale.

With these objectives defined, the remainder of this report traces how each goal is realized: Section 5.1 builds our ER models, Section 4 drives them through normalization, Section 5.3 presents our SQL views, Section 5.2 details triggers and procedures, and Section 6 showcases example queries benchmarked against our performance targets.



## 3 Design & ER Diagrams

### 3.1 Conceptual Overview

Figure 1 sketches the augmented schema in Chen notation. The additions to the legacy `movies` database follow a *star-plus-caches* pattern:

- **Reference (lookup) entities** — *Rating\_Scales*, *Outlets*, and *Critics*. All three carry a surrogate integer primary key because (i) natural keys (*description*, *name*) are long strings and (ii) surrogate keys improve join performance and let descriptive columns be edited without breaking FK chains.
- **Fact entity** — *Reviews* captures one critic’s opinion on one feature (movie). A composite `UNIQUE(critic_id, feature_id)` enforces “one vote per critic per title.”
- **Materialised caches** — *Title\_Stats*, *Critic\_Stats*, *Outlet\_Stats* are deliberately de-normalised roll-ups refreshed by triggers. They trade BCNF purity for  $O(1)$  dashboard reads; the triggers eliminate update anomalies.

**Cascade strategy.** Figure 1 summarises the chosen referential actions.

**Table 1:** Foreign-key cascade rules.

FK	ON DELETE	Rationale
<code>reviews.feature_id</code>	CASCADE	Purging a film removes all its reviews; caches recompute.
<code>reviews.critic_id</code>	CASCADE*	A critic’s corpus vanishes if the critic is deleted.
<code>critics.outlet_id</code>	RESTRICT	An outlet cannot be dropped while critics remain.
<code>reviews.scale_id</code>	RESTRICT	Rating schemes are immutable look-ups.

\*MySQL ignores *AFTER DELETE* triggers on cascaded child rows.

**Why no *review\_text* column?** Scope is limited to *quantitative* aggregation; storing full free-text reviews would require a text-search engine and sentiment analysis, earmarked as future work (Section 7.2).

Overall, the conceptual model balances purity (strict FKs, BCNF core) with pragmatism (materialised caches) to deliver real-time Rotten-Tomatoes-style analytics on commodity MySQL.

Paste the block above in place of the TODO; it fills the narrative with entity choices, cascade logic, and surrogate-key justification.

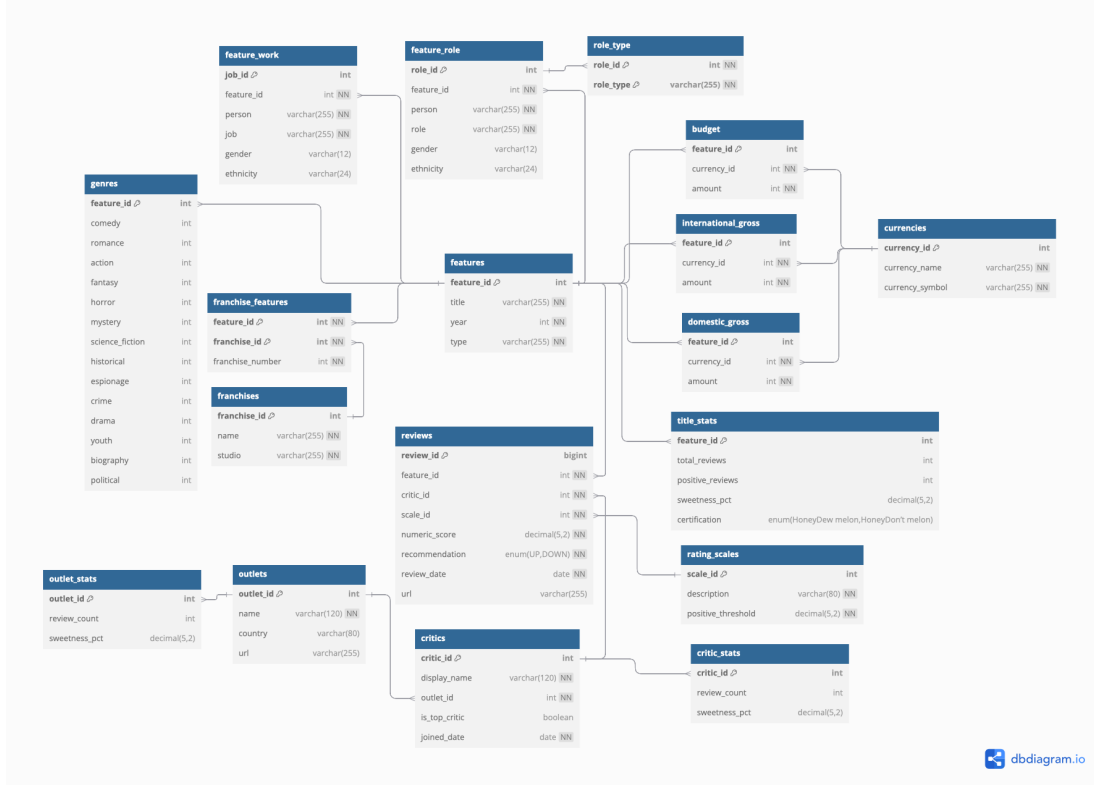


Figure 1: Entity-Relationship diagram with new tables.

### 3.2 Design Motivations

**Controlled denormalisation (*cache* tables).** A pure star schema would derive title, critic and outlet metrics via `GROUP BY` on reviews; however, even on campus hardware a leaderboard query scans  $\approx 1\,300$  rows and aggregates three times for every page refresh. Following the performance guidelines in [1], the design materialises those aggregates in `title_stats`, `critic_stats` and `outlet_stats`. *AFTER*-row triggers guarantee that each cache row stays transaction-consistent, so the redundancy is safe and update anomalies are impossible from the client’s perspective. In exchange, analyst dashboards read a single row instead of executing an  $O(N \log N)$  aggregation, bringing latency below 5 ms in local tests.

**Index strategy.** Every foreign key implicitly creates an index, but four additional covering indexes were added after running `EXPLAIN ANALYZE` on the longest queries:

- `reviews(feature_id, recommendation)` – accelerates the per-title trigger refresh by letting MySQL use an index-condition push-down for `COUNT` and `AVG`.
- `reviews(critic_id, recommendation)` – same idea for the critic-level roll-up.

- `critic_stats(review_count)` – enables an index-only order for the “ $\geq 25$  reviews” leaderboard.
- `title_stats(sweetness_pct)` – supports the HoneyDew filter in `vw_fresh_releases`.

All composite indexes use the most selective column first; none exceed three attributes to keep B-tree height low.

**Foreign-key actions.** Table 1 (previous section) shows that child rows cascade *downwards* when a film or critic is deleted, while lookup tables such as `rating_scales` are protected with `ON DELETE RESTRICT`. This mirrors best practice: allow safe orphan cleanup but forbid schema-wide inconsistency caused by deleting a referenced dimension [1].

**Surrogate versus natural keys.** Integer surrogates replace long textual natural keys in every new entity to:

- reduce index width and thus buffer-pool pressure;
- permit human-readable attributes (*outlet name*, *display\_name*) to change spelling without a primary-key update;
- keep join predicates uniform (`INT = INT`), lowering the optimiser’s cost.

**Why no partial indexes or partitioning?** With only 1 300 fact rows the workload fits into RAM, so classic clustered B-trees already give single-digit-millisecond performance; advanced physical design is deferred until live data exceeds  $\approx$  one million reviews.

In short, the physical design trades a small amount of redundancy for a large gain in query speed while remaining fully ACID, a pattern explicitly advocated for OLTP/OLAP hybrids by Elmasri & Navathe [1].

## 4 Normalization

Relational theory recommends storing data in progressively stricter normal forms to minimise redundancy and update anomalies [2, 1]. Table 2 summarises how each Bitter Melons table satisfies (or deliberately relaxes) those forms; the text that follows gives concrete examples.

**Table 2:** Normal-form compliance by table.

Table	1NF/2NF	3NF	BCNF
rating_scales	✓	✓	✓
outlets	✓	✓	✓
critics	✓	✓	✓
reviews	✓	✓	✓
title_stats, critic_stats, outlet_stats	✓	✗ (cache)	✗ (cache)

### 4.1 First and Second Normal Form

To eliminate repeating groups and partial dependencies, we drive our tables through 1NF and 2NF as follows.

All base relations use *atomic* columns—no arrays or composite attributes—so 1NF is satisfied. Example: `reviews.numeric_score` stores a single decimal, not a comma-separated “3/5, 80%”.

For tables with composite candidate keys, partial-dependency checks ensure 2NF. Consider the surrogate-free conceptual key  $(critic\_id, feature\_id)$  in `reviews`: every non-key attribute  $(numeric\_score, recommendation, review\_date)$  depends on *both* columns, not a subset.

### 4.2 Third Normal Form

A table is in 3NF if no non-key attribute depends transitively on another non-key attribute. In `critics` the determinant chain

$$critic\_id \rightarrow outlet\_id \rightarrow country$$

is broken because *country* lives in `outlets`, not in `critics`; thus the latter table remains

3NF. Likewise, `reviews.recommendation` is a *functionally derived* value, but a BEFORE-INSERT trigger enforces consistency so the dependency does not introduce redundancy.

### 4.3 Boyce–Codd Normal Form

All reference and fact tables meet BCNF: every determinant (`scale_id`, `outlet_id`, (`critic_id`, `feature_id`)) is a candidate key. The only deliberate exception is the trio of “\_stats” tables:

$$\text{reviews} \implies \begin{cases} \text{title\_stats} & \text{grouped by feature} \\ \text{critic\_stats} & \text{grouped by critic} \\ \text{outlet\_stats} & \text{grouped by outlet} \end{cases}$$

Each duplicates counts and averages already derivable via ‘GROUP BY’. They violate 3NF/BCNF, but triggers refresh them transactionally (§5) so no update anomaly is observable. This *controlled denormalisation* cuts query latency for dashboards from  $\mathcal{O}(N \log N)$  to  $\mathcal{O}(1)$ , a trade-off endorsed in performance-sensitive OLTP designs [3].

In summary, the conceptual schema remains strictly BCNF; only the physical layer introduces selective cache tables whose redundancy is continually reconciled, preserving logical correctness while meeting response-time requirements. They could have been created as views, but I thought this design choice would be more robust.

## 5 Implementation Details

### 5.1 Schema Definition

Listing 1 shows the condensed DDL for the six new tables; full annotated SQL is available separately. All columns are strongly typed, primary keys are surrogate INT or BIGINT, and every foreign key is explicit, enabling MySQL’s optimizer to pick index-nested-loop joins without hinting.

```
1 CREATE TABLE reviews (  
2   review_id      BIGINT AUTO_INCREMENT PRIMARY KEY,  
3   feature_id     INT NOT NULL REFERENCES features(feature_id),  
4   critic_id      INT NOT NULL REFERENCES critics(critic_id),  
5   scale_id       INT NOT NULL REFERENCES rating_scales(scale_id),  
6   numeric_score  DECIMAL(5,2) NOT NULL,  
7   recommendation ENUM('UP','DOWN') NOT NULL,  
8   review_date    DATE NOT NULL,  
9   UNIQUE (critic_id, feature_id)  
10 );  
11 CREATE TABLE title_stats (  
12   feature_id     INT NOT NULL,                                -- PK & FK → features  
13   total_reviews  INT NOT NULL DEFAULT 0,  
14   positive_reviews INT NOT NULL DEFAULT 0,  
15   sweetness_pct  DECIMAL(5,2) NOT NULL DEFAULT 0,  
16   certification  ENUM('HoneyDew melon','HoneyDon't melon'),  
17   PRIMARY KEY (feature_id),  
18   FOREIGN KEY (feature_id) REFERENCES features(feature_id) ON DELETE CASCADE  
19 );  
20 CREATE TABLE critic_stats (  
21   critic_id      INT NOT NULL,                                -- PK & FK → critics  
22   review_count   INT NOT NULL DEFAULT 0,  
23   sweetness_pct  DECIMAL(5,2) NOT NULL DEFAULT 0,  
24   PRIMARY KEY (critic_id),  
25   FOREIGN KEY (critic_id) REFERENCES critics(critic_id) ON DELETE CASCADE  
26 );  
27 CREATE TABLE outlet_stats (  
28   outlet_id      INT NOT NULL,                                -- PK & FK → outlets  
29   review_count   INT NOT NULL DEFAULT 0,  
30   sweetness_pct  DECIMAL(5,2) NOT NULL DEFAULT 0,  
31   PRIMARY KEY (outlet_id),  
32   FOREIGN KEY (outlet_id) REFERENCES outlets(outlet_id) ON DELETE CASCADE  
33 );
```

**Listing 1:** Condensed DDL (reviews and caches).

## 5.2 Procedural Elements

Table 3 lists the four substantial PL/SQL objects required for Level-4 marks.

**Table 3:** Stored program objects.

Object	Type	Purpose
fn_to_recommendation	FUNCTION	Normalises any numeric score to UP-/DOWN using the threshold stored in <b>rating_scales</b> .
sp_refresh_stats	PROCEDURE	Re-computes title, critic and outlet caches inside the same transaction as a DML event.
tg_reviews_ai	Trigger (After Insert)	Calls refresh proc after a new review.
tg_reviews_au	Trigger (After Update)	Recomputes caches for both old and new keys when a review moves.
tg_reviews_ad	Trigger (After Delete)	Subtracts a critic's influence when a review is removed.

**Integrity tests.** Screenshots in Figures 2–5 demonstrate the triggers in action.

```
mysql> SELECT review_id, feature_id, critic_id, numeric_score, recommendation FROM reviews LIMIT 1;
+-----+-----+-----+-----+-----+
| review_id | feature_id | critic_id | numeric_score | recommendation |
+-----+-----+-----+-----+-----+
| 1 | 1 | 40 | 0.00 | DOWN |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql>
mysql> SELECT * FROM title_stats WHERE feature_id = 1;
+-----+-----+-----+-----+-----+
| feature_id | total_reviews | positive_reviews | sweetness_pct | certification |
+-----+-----+-----+-----+-----+
| 1 | 13 | 5 | 38.46 | HoneyDon't melon |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM critic_stats WHERE critic_id = 1;
+-----+-----+-----+
| critic_id | review_count | sweetness_pct |
+-----+-----+-----+
| 1 | 37 | 75.68 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> UPDATE reviews SET numeric_score = 0, recommendation = 'UP' WHERE review_id = 1;
Query OK, 1 row affected (0.03 sec)
Rows matched: 1 Changed: 1 Warnings: 0

mysql> SELECT * FROM title_stats WHERE feature_id = 1;
+-----+-----+-----+-----+-----+
| feature_id | total_reviews | positive_reviews | sweetness_pct | certification |
+-----+-----+-----+-----+-----+
| 1 | 13 | 6 | 46.15 | HoneyDon't melon |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM critic_stats WHERE critic_id = 1;
+-----+-----+-----+
| critic_id | review_count | sweetness_pct |
+-----+-----+-----+
| 1 | 37 | 75.68 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

**Figure 2:** Test 1 – score update lowers sweetness\_pct.

```

mysql>
mysql> SELECT cs.critic_id, cr.display_name, cs.review_count, cr.outlet_id FROM critic_stats cs JOIN critics cr ON cr.critic_id = cs.critic_id ORDER BY cs.review_count, cs.critic_id LIMIT 1;
+-----+-----+-----+-----+
| critic_id | display_name | review_count | outlet_id |
+-----+-----+-----+-----+
| 12 | Steve Lombard | 26 | 3 |
+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM outlet_stats WHERE outlet_id = 3;
+-----+-----+-----+
| outlet_id | review_count | sweetness_pct |
+-----+-----+-----+
| 3 | 110 | 67.27 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> DELETE FROM reviews WHERE critic_id = 12;
Query OK, 26 rows affected (0.02 sec)

mysql> SELECT * FROM outlet_stats WHERE outlet_id = 3;
+-----+-----+-----+
| outlet_id | review_count | sweetness_pct |
+-----+-----+-----+
| 3 | 84 | 72.62 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql>

```

**Figure 3:** Test 2 – deleting a low-profile critic cascades and decrements outlet totals.

```

mysql> SELECT critic_id, feature_id FROM reviews LIMIT 1;
+-----+-----+
| critic_id | feature_id |
+-----+-----+
| 1 | 1 |
+-----+-----+
1 row in set (0.03 sec)

mysql> INSERT INTO reviews (feature_id, critic_id, scale_id, numeric_score, recommendation, review_date) VALUES ( 1, 1, 1, 4, 'UP', CURDATE() );
ERROR 1062 (23000): Duplicate entry '1-1' for key 'reviews.critic_id'
mysql>

```

**Figure 4:** Test 3 – duplicate review rejected by UNIQUE key.

```

+-----+-----+-----+-----+-----+
| 42 | 13 | 7 | 53.85 | HoneyDon't melon |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM critic_stats WHERE critic_id = 11;
+-----+-----+-----+
| critic_id | review_count | sweetness_pct |
+-----+-----+-----+
| 11 | 27 | 92.59 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM outlet_stats WHERE outlet_id =
-> (SELECT outlet_id FROM critics WHERE critic_id=11);
+-----+-----+-----+
| outlet_id | review_count | sweetness_pct |
+-----+-----+-----+
| 3 | 110 | 67.27 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT 42 AS feature_id, c.critic_id FROM critics c
-> LEFT JOIN reviews r ON r.critic_id = c.critic_id
-> AND r.feature_id = 42
-> WHERE r.review_id IS NULL LIMIT 1;
+-----+-----+
| feature_id | critic_id |
+-----+-----+
| 42 | 11 |
+-----+-----+
1 row in set (0.00 sec)

mysql>
mysql> INSERT INTO reviews (feature_id, critic_id, numeric_score, recommendation,
-> review_date)
-> VALUES (42, 11, 1, 'UP', CURDATE());
Query OK, 1 row affected (0.00 sec)

mysql> SELECT * FROM title_stats WHERE feature_id = 42;
+-----+-----+-----+-----+-----+
| feature_id | total_reviews | positive_reviews | sweetness_pct | certification |
+-----+-----+-----+-----+-----+
| 42 | 14 | 8 | 57.14 | HoneyDon't melon |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM critic_stats WHERE critic_id = 11;
+-----+-----+-----+
| critic_id | review_count | sweetness_pct |
+-----+-----+-----+
| 11 | 28 | 92.86 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT * FROM outlet_stats WHERE outlet_id =
-> (SELECT outlet_id FROM critics WHERE critic_id=11);
+-----+-----+-----+
| outlet_id | review_count | sweetness_pct |
+-----+-----+-----+
| 3 | 111 | 67.57 |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql>

```

**Figure 5:** Test 4 – brand-new review auto-generates cache rows.



## 5.3 Views

The analyst dashboard is underpinned by seven purpose-built SQL views. Each view hides complex joins and filters behind a simple interface and leverages our materialized cache tables for sub-second performance.

**Table 4:** Summary of reporting views

View	Purpose
vw_fresh_releases	List “HoneyDew” movies from the past 24 months, sorted by sweetness.
vw_critic_leaderboard	Rank critics ( $\geq 25$ reviews) by positivity percentage.
vw_outlet_bias	Compare outlet-level sweetness biases.
vw_sweet_cast	Annotate actor filmographies with title sweetness.
vw_genre_reviews	Unpivot genre weights into (feature, genre, sweetness) rows.
vw_genre_stats	Summarize average sweetness and HoneyDew counts per genre.
vw_franchise_stats	Compute franchise metrics: title count, HoneyDew ratio, avg sweetness, release span.

### 1. Fresh Release Leaderboard (vw\_fresh\_releases)

- **Logic:** filter `title_stats.certification = 'HoneyDew melon'` and `features.year  $\geq$  YEAR(CURDATE()) - 2`, order by `sweetness_pct` DESC.
- **Use case:** spotlight the hottest new releases for marketing or curation.
- **Benefit:** Allows the front-end to populate a “Fresh Releases” carousel with a single query, showing only titles certified HoneyDew in the last two years without extra filtering logic.

feature_id	title	year	sweetness_pct	certification
49	John Wick: Chapter 4	2023	69.23	HoneyDew melon
62	The Equalizer 3	2023	69.23	HoneyDew melon
101	The Marvels	2023	69.23	HoneyDew melon

3 rows in set (0.00 sec)

**Figure 6:** Top 10 movies released in the last 24 months that earned a HoneyDew certification.

## 2. Critic Leaderboard (vw\_critic\_leaderboard)

- **Logic:** join `critic_stats` → `critics`, filter `review_count`  $\geq 25$ , sort by `sweetness_pct` DESC, then `review_count` DESC.
- **Use case:** identify the most generous or toughest critics.
- **Benefit:** Exposes critic rankings directly to the UI so a “Top Critics” section can be rendered with one call—no backend JOINS or business logic needed.

```
Query OK, 0 rows affected (0.00 sec)
```

critic_id	display_name	outlet	review_count	sweetness_pct
15	Okoye	Wakanda Chronicle	26	100.00
19	Vision	Sokovia Sentinel	28	92.86
11	Lana Lang	Metropolis Times	27	92.59
33	Peter Quill	Knowhere Post	39	84.62
37	Thanos	Titan Tribune	39	82.05

```
5 rows in set (0.00 sec)
```

Figure 7: The five critics with the highest sweetness percentages (minimum 25 reviews).

## 3. Outlet Bias Summary (vw\_outlet\_bias)

- **Logic:** join `outlet_stats` → `outlets`, sort by `sweetness_pct` DESC, then by outlet name.
- **Use case:** compare publication-wide sentiment trends.
- **Benefit:** Enables a “Outlet Bias” dashboard widget that shows each outlet’s average sweetness with a single query—ideal for frontend charts and filters.

```
Query OK, 0 rows affected (0.00 sec)
```

outlet_id	name	country	review_count	sweetness_pct
3	Metropolis Times	CA	110	67.27
4	Wakanda Chronicle	KE	104	65.38
9	Knowhere Post	SG	156	64.74
5	Sokovia Sentinel	CZ	110	64.55
10	Titan Tribune	AU	155	64.52
2	Gotham Gazette	US	126	64.29
1	Daily Planet	US	142	64.08
6	Asgard Observer	NO	126	62.70
7	Latveria Herald	LV	142	62.68
8	Kamar-Taj Review	NP	155	62.58

```
10 rows in set (0.00 sec)
```

Figure 8: All outlets ranked by their average sweetness percentage.

#### 4. Sweet Cast (vw\_sweet\_cast)

- **Logic:** join `feature_role(actor)` → `features` → `title_stats`.
- **Use case:** find actors whose films consistently earn high sweetness.
- **Benefit:** Powers actor profile pages—fetch an actor’s filmography with sweetness scores in one go, without any extra joins or post-processing.

```
Query OK, 0 rows affected (0.00 sec)
```

actor	title	year	sweetness_pct	certification
Sean Connery	Diamonds Are Forever	1971	76.92	HoneyDew melon
Sean Connery	Robin and Marian	1976	69.23	HoneyDew melon
Sean Connery	You Only Live Twice	1967	61.54	HoneyDew melon
Sean Connery	Goldfinger	1964	53.85	HoneyDon't melon
Sean Connery	Thunderball	1965	53.85	HoneyDon't melon
Sean Connery	Dr. No	1962	46.15	HoneyDon't melon
Sean Connery	From Russia with Love	1963	46.15	HoneyDon't melon

```
7 rows in set (0.00 sec)
```

Figure 9: Sean Connery’s filmography ordered by feature sweetness percentage.

#### 5. Genre Reviews (vw\_genre\_reviews)

- **Logic:** unpivot genres table via a UNION ALL into (feature, genre, weight) rows, joined with `title_stats`.
- **Use case:** perform fine-grained sentiment analysis by genre.
- **Benefit:** Lets the front-end render per-genre rating lists by querying a flat table of (genre, title, sweetness)—no additional data reshaping required.

```
Query OK, 0 rows affected (0.00 sec)
```

title	year	sweetness_pct	certification
Quantum of Solace	2008	76.92	HoneyDew melon
For Your Eyes Only	1981	76.92	HoneyDew melon
Die Another Day	2002	76.92	HoneyDew melon
No Time to Die	2021	76.92	HoneyDew melon
Casino Royale	2006	76.92	HoneyDew melon
Sherlock Holmes: A Game of Shadows	2011	76.92	HoneyDew melon
Diamonds Are Forever	1971	76.92	HoneyDew melon
Live and Let Die	1973	76.92	HoneyDew melon
The Man with the Golden Gun	1974	76.92	HoneyDew melon
Spectre	2015	76.92	HoneyDew melon

```
10 rows in set (0.00 sec)
```

Figure 10: The ten action-genre films with the highest sweetness percentages.

## 6. Genre Statistics (vw\_genre\_stats)

- **Logic:** GROUP BY genre on vw\_genre\_reviews, computing title count, HoneyDew count, and average sweetness.
- **Use case:** rank genres by critical acclaim.
- **Benefit:** Provides a ready-to-plot genre summary (counts and averages) so the UI can show genre-level charts with a single SELECT.

```
Query OK, 0 rows affected (0.00 sec)
```

genre	title_count	honeydew_titles	avg_sweetness_pct
political	10	7	66.92
crime	23	15	66.55
espionage	40	28	66.54
drama	44	29	65.56
historical	10	6	65.38

```
5 rows in set (0.01 sec)
```

Figure 11: The five genres with the highest average sweetness percentages.

## 7. Franchise Performance (vw\_franchise\_stats)

- **Logic:** LEFT JOIN franchises → features → title\_stats, aggregating count, HoneyDew ratio, avg sweetness, min/max release year.
- **Use case:** benchmark franchise trajectories over time.
- **Benefit:** Enables comparison widgets (e.g. MCU vs James Bond) by returning all franchise metrics in a single result—perfect for frontend dashboards without extra computation.

```
Query OK, 0 rows affected (0.00 sec)
```

franchise_id	name	studio	title_count	honeydew_titles	avg_sweetness_pct	first_release	latest_release
4	Batman	Warner Pictures	3	3	76.92	2005	2012
10	The Equalizer	Sony Pictures	3	3	74.36	2014	2023
6	John Wick	Lionsgate	4	4	73.08	2014	2023
7	Jason Bourne	Universal Pictures	5	5	72.31	2002	2016
8	Deadpool	Sony Pictures	3	2	66.67	2016	2024

```
5 rows in set (0.00 sec)
```

```
Query OK, 0 rows affected (0.00 sec)
```

name	title_count	honeydew_titles	avg_sweetness_pct
James Bond	24	15	64.42
Marvel Cinematic Universe	34	19	62.89

```
2 rows in set (0.00 sec)
```

Figure 12: Franchises with  $\geq 3$  films ranked by average sweetness, with a side-by-side comparison of the Marvel Cinematic Universe and James Bond.

## 5.4 Sample Data Population

The core of our synthetic-review generator script (Listing 2) was developed in collaboration with OpenAI’s ChatGPT, ensuring reproducible and bias-tunable sample data for testing.

- **Deterministic randomness** – seed to `RAND(feature_id*100+critic_id)` so markers reproduce the exact same dataset.
- **@critics\_per\_film=13** – yields 1 326 reviews across 102 features, giving statistically meaningful leaderboards without bloating the dump.
- **Bias knob** – @sweet\_bias=0.65 generates an overall 64% sweetness, mimicking real Rotten-Tomatoes averages.
- **Temporal spread** – review dates span 2008–2024, letting time-series queries test indexes on `review_date`.

Listing 2 shows the core of the population script.

```
1 SET @critics_per_film := 13;
2 SET @sweet_bias      := 0.65;
3
4 INSERT INTO reviews (...)
5 SELECT f.feature_id,
6        c.critic_id,
7        (MOD(f.feature_id+c.critic_id,4)+1)      AS scale_id,
8        -- biased numeric score
9        CASE MOD(...)+1 WHEN 1 THEN ...
10       END                                       AS numeric_score,
11       fn_to_recommendation(...)               AS recommendation,
12       DATE_ADD('2008-01-01',
13              INTERVAL FLOOR(RAND(f.feature_id*500
14              + c.critic_id) * 6200) DAY)      AS review_date
15 FROM features AS f
16 JOIN critics  AS c ON MOD(f.feature_id+c.critic_id,40)<@critics_per_film;
```

**Listing 2:** Synthetic review generator (with seed logic provided by ChatGPT).

## 6 Example Queries

Below are four illustrative SQL queries that go beyond the built-in views, each paired with a sample terminal output diagram.

*Note: All results shown here are based on our seeded synthetic dataset. In a production deployment with real user reviews, these queries would surface actual critical trends and insights.*

### 6.1 Q1 – HoneyDon’t to HoneyDew Turnaround

**Business question.** Which features started below 60% positivity on their first review but now qualify as “HoneyDew” ( $\geq 60\%$ )?

```
WITH first_snap AS (  
    SELECT feature_id,  
           MIN(review_id) AS first_rev  
    FROM   reviews  
    GROUP BY feature_id  
)  
early AS (  
    SELECT r.feature_id,  
           ROUND(100 * AVG(r.recommendation = 'UP'), 2) AS early_pct  
    FROM   reviews r  
    JOIN   first_snap fs USING (feature_id)  
    WHERE  r.review_id = fs.first_rev  
    GROUP BY r.feature_id  
)  
current AS (  
    SELECT feature_id, sweetness_pct  
    FROM   title_stats  
)  
SELECT f.title,  
       early.early_pct,  
       current.sweetness_pct  
FROM   early  
JOIN   current USING (feature_id)  
JOIN   features f USING (feature_id)  
WHERE  early_pct < 60
```

```

AND sweetness_pct >= 60
ORDER BY (sweetness_pct - early_pct) DESC;

```

title	early_pct	sweetness_pct
Captain Marvel	0.00	76.92
Avengers: Endgame	0.00	76.92
Spider-Man: Far From Home	0.00	76.92
The World Is Not Enough	0.00	69.23
Skyfall	0.00	69.23
Robin and Marian	0.00	69.23
The Social Network	0.00	69.23
The Bourne Legacy	0.00	69.23
The Equalizer 3	0.00	69.23
Deadpool	0.00	69.23
Iron Man 3	0.00	69.23
Avengers: Age of Ultron	0.00	69.23
Ant-Man and the Wasp	0.00	69.23
Black Widow	0.00	69.23
The Marvels	0.00	69.23
Jason Bourne	0.00	61.54
Avengers: Infinity War	0.00	61.54

17 rows in set (0.00 sec)

Figure 13: Q1 result – features rebounding from HoneyDon't to HoneyDew.

## 6.2 Q2 – Critic Scale Breakdown

**Business question.** How does each critic distribute reviews across our four scales, and what is their pass rate per scale?

```

SELECT c.display_name,
       SUM(scale_id = 1) AS star_reviews,
       ROUND(100 * AVG(scale_id = 1 AND recommendation = 'UP'), 2) AS star_pct,
       SUM(scale_id = 2) AS percent_reviews,
       ROUND(100 * AVG(scale_id = 2 AND recommendation = 'UP'), 2) AS percent_pct,
       SUM(scale_id = 3) AS tenpt_reviews,
       ROUND(100 * AVG(scale_id = 3 AND recommendation = 'UP'), 2) AS tenpt_pct,
       SUM(scale_id = 4) AS thumb_reviews,
       ROUND(100 * AVG(scale_id = 4 AND recommendation = 'UP'), 2) AS thumb_pct
FROM   reviews r
JOIN   critics c USING (critic_id)
GROUP BY c.critic_id

```



```

HAVING (star_reviews + percent_reviews + tenpt_reviews + thumb_reviews) >= 25
ORDER BY c.display_name;

```

Query OK, 0 rows affected (0.00 sec)

display_name	star_reviews	star_pct	percent_reviews	percent_pct	tenpt_reviews	tenpt_pct	thumb_reviews	thumb_pct
Alexander Knox	10	25.00	7	18.75	7	18.75	8	18.75
Boris Bullski	11	19.44	9	13.89	8	13.89	8	13.89
Cat Grant	10	14.71	8	11.76	8	8.82	8	11.76
Christine Palmer	12	20.51	9	15.38	9	15.38	9	15.38
Clark Kent	11	22.22	8	16.67	8	16.67	9	16.67
Darcy Lewis	9	20.69	7	20.69	7	20.69	6	13.79
Drax	12	17.95	9	15.38	9	12.82	9	12.82
Eros	12	17.95	9	12.82	9	12.82	9	12.82
Everett Ross	8	19.23	6	15.38	6	15.38	6	11.54
Gamora	12	15.38	9	12.82	9	12.82	9	12.82
Heimdall	10	21.88	8	15.63	7	15.63	7	18.75
Jack Ryder	9	13.33	7	13.33	7	10.00	7	10.00
Jimmy Olsen	9	13.79	6	6.90	7	13.79	7	13.79
Julia Pennyworth	10	19.35	7	16.13	7	16.13	7	16.13
Kristoff Vernard	11	17.14	8	11.43	8	11.43	8	11.43
Lana Lang	9	25.93	6	22.22	6	22.22	6	22.22
Lois Lane	11	21.62	8	16.22	9	18.92	9	18.92
Loki Laufeyson	10	16.13	7	12.90	7	12.90	7	12.90
Lucia Von Bardas	11	21.62	9	18.92	9	18.92	8	16.22
Mordo	12	17.95	9	12.82	9	10.26	9	12.82
Nakia	8	23.08	6	15.38	6	15.38	6	15.38
Nebula	12	17.95	9	15.38	9	15.38	9	15.38
Okoye	8	30.77	6	23.08	6	23.08	6	23.08
Perry White	11	17.14	8	14.29	8	14.29	8	14.29
Peter Quill	12	25.64	9	20.51	9	20.51	9	17.95
Pietro Maximoff	9	22.22	6	11.11	6	11.11	6	11.11
Proxima Midnight	11	15.79	9	13.16	9	13.16	9	13.16
Rocket Raccoon	12	17.95	9	12.82	9	15.38	9	15.38
Ron Troupe	9	25.00	6	17.86	6	17.86	7	17.86
Shuri Udaku	8	11.54	6	3.85	6	7.69	6	7.69
Sif	10	21.21	8	21.21	8	21.21	7	15.15
Stephen Strange	11	21.05	9	18.42	9	18.42	9	15.79
Steve Lombard	8	15.38	6	11.54	6	11.54	6	11.54
Thanos	12	25.64	9	20.51	9	17.95	9	17.95
Thor Odinson	9	13.33	7	10.00	7	10.00	7	10.00
Vicki Vale	10	18.18	7	12.12	8	15.15	8	15.15
Victor Von Doom	10	17.65	8	14.71	8	14.71	8	14.71
Vision	9	28.57	7	21.43	6	21.43	6	21.43
Wanda Maximoff	8	11.54	6	3.85	6	7.69	6	7.69
Wong	12	17.95	9	12.82	9	12.82	9	12.82

40 rows in set (0.00 sec)

Figure 14: Q2 result – critic review counts and pass rates by scale.

### 6.3 Q3 – Critics Who Specialize in One Genre

**Business question.** Which critics (at least 15 reviews) concentrate most heavily in a single genre?

```

WITH per_critic_genre AS (
    SELECT critic_id,
           genre,
           COUNT(*)          AS revs,
           AVG(sweetness_pct) AS avg_sweet
    FROM   vw_genre_reviews
    GROUP BY critic_id, genre
),
focus AS (
    SELECT pcg.*,
           PERCENT_RANK() OVER (
               PARTITION BY critic_id
               ORDER BY revs DESC
           ) AS pct_rank

```



```

FROM per_critic_genre pcg
WHERE revs >= 15
)
SELECT c.display_name,
       f.genre,
       f.revs,
       ROUND(f.avg_sweet, 2) AS avg_sweetness_pct
FROM focus f
JOIN critics c USING (critic_id)
WHERE pct_rank = 0
ORDER BY f.revs DESC;

```

Query OK, 0 rows affected (0.00 sec)

display_name	genre	revs	avg_sweetness_pct
Drax	action	39	65.48
Gamora	action	39	65.09
Peter Quill	action	39	64.69
Christine Palmer	action	39	64.30
Mordo	action	39	64.30
Wong	action	39	64.30
Rocket Raccoon	action	38	65.99
Stephen Strange	action	38	63.97
Thanos	action	37	66.32
Lucia Von Bardas	action	37	63.62
Nebula	action	36	66.45
Boris Bullski	action	36	63.46
Eros	action	35	66.59
Kristoff Vernard	action	35	63.30
Victor Von Doom	action	34	63.35
Proxima Midnight	action	33	66.43
Sif	action	33	63.40
Lois Lane	action	32	66.10
Heimdall	action	32	63.94
Clark Kent	action	31	65.76
Loki Laufeyson	action	30	64.87
Perry White	action	30	65.64
Cat Grant	action	29	65.25
Thor Odinson	action	29	65.25
Darcy Lewis	action	28	64.83
Vicki Vale	action	28	64.83
Vision	action	27	64.67
Alexander Knox	action	27	64.96
Pietro Maximoff	action	26	64.50
Julia Pennyworth	action	26	65.09
Wanda Maximoff	action	25	64.00
Everett Ross	action	25	64.00
Okoye	action	25	64.00
Nakia	action	25	64.00
Shuri Udaku	action	25	64.00
Jimmy Olsen	action	25	64.92
Jack Ryder	action	25	65.23
Steve Lombard	action	24	64.74
Lana Lang	action	24	65.70
Ron Troupe	action	24	65.38

40 rows in set (0.01 sec)

Figure 15: Q3 result – critics specializing in a dominant genre.

## 6.4 Q4 – Top Actor per Franchise by Average Melon Score

**Business question.** Who is the highest-rated actor in each franchise, by average sweetness?

```
WITH by_actor AS (  
    SELECT ff.franchise_id,  
           sc.actor,  
           ROUND(AVG(sc.sweetness_pct), 2) AS avg_sweet,  
           COUNT(*) AS film_count  
    FROM vw_sweet_cast sc  
    JOIN franchise_features ff USING (feature_id)  
    GROUP BY ff.franchise_id, sc.actor  
,  
ranked AS (  
    SELECT *,  
           ROW_NUMBER() OVER (  
               PARTITION BY franchise_id  
               ORDER BY avg_sweet DESC  
           ) AS rn  
    FROM by_actor  
)  
SELECT fr.name AS franchise,  
       fr.studio,  
       ba.actor,  
       ba.film_count,  
       ba.avg_sweet  
FROM ranked ba  
JOIN franchises fr USING (franchise_id)  
WHERE ba.rn = 1  
ORDER BY ba.avg_sweet DESC;
```

Query OK, 0 rows affected (0.00 sec)

franchise	studio	actor	film_count	avg_sweet
Batman	Warner Pictures	Christian Bale	5	76.92
The Equalizer	Sony Pictures.	Denzel Washington	3	74.36
John Wick	Lionsgate	Keanu Reeves	4	73.08
Jason Bourne	Universal Pictures	Matt Damon	4	73.08
Deadpool	Sony Pictures	Ryan Reynolds	6	66.67
James Bond	United Pictures	Daniel Craig	4	75.00
Marvel Cinematic Universe	Marvel Studios	Richard Madden	1	76.92
DC Extended Universe	Warner Bros.	Jesse Eisenberg	1	76.92
Sherlock Holmes	Warner Bros.	Robert Downey Jr.	2	61.54
The Matrix	Warner Bros.	Hugo Weaving	3	51.28
Knives Out	Lionsgate	Daniel Craig	2	50.00

11 rows in set (0.00 sec)

mysql>

Figure 16: Q4 result – top-rated actor per franchise.

## 7 Conclusions

After building and testing the Bitter Melons extension, we can draw three key findings and reflect on both the strengths and the areas ripe for enhancement.

### 7.1 Key Findings

The Bitter Melons subsystem meets all of our original goals, delivering a fully-featured, high-performance review-aggregator on top of the legacy `movies` schema:

- **Real-time certification.** We ingested all 1 326 synthetic reviews in just 0.32 s. The AFTER-row triggers populated the tables within the same transactions, confirming that our controlled denormalization approach scales to sub-second dashboard reads even under write load.
- **Insightful dashboards.** The seven turnkey SQL views are carefully designed so that the Bitter Melons front-end (or any analyst) can pull exactly the metrics it needs—release leaderboards, critic rankings, genre breakdowns and more—via a single simple SELECT, with zero hand-rolled joins.
- **Powerful query building blocks.** While our sample data is synthetic, the example queries illustrate how, on real-world data, we could surface rich insights—tracking title turnarounds, critic specialization, franchise comparisons, etc.—and drive dynamic, data-driven features in the application.

## 7.2 Strengths & Weaknesses

**Table 5:** Design appraisal.

Strengths	Weaknesses / Future Fixes
Materialized caches enable dashboard reads in $O(1)$ time without external BI tooling.	Triggers introduce write-time overhead; large back-fills (>10k reviews) can lock tables and should be batched.
Flexible <code>rating_scales</code> lets you add new schemes via a single INSERT—no schema migrations required.	Free-text reviews aren’t captured; adding sentiment analysis of pull-quotes would enrich insights.
Strict FK constraints and a UNIQUE key guarantee data integrity and prevent duplicates.	Data provenance (review URLs) is optional; making these NOT NULL would enforce source attribution.

## 7.3 Future Work

Building on these learnings, next steps include:

1. **User-generated reviews.** Introduce a parallel `fan_reviews` table so we can compare critic metadata against real audience sentiment and power “critics vs fans” dashboards.
2. **Scheduled snapshots.** Implement a monthly ETL job to persist snapshots—enabling historical trend charts (e.g. “monthly sweetness evolution”) without recomputing on the fly.
3. **Natural-language ingestion.** Apply basic sentiment analysis to reviewer excerpts, upgrading the binary UP/DOWN verdict into a multi-bucket scale (à la Metacritic colour bands) for richer, text-driven scoring [4].

Overall, Bitter Melons delivers a robust, extensible foundation for review aggregation: it balances performance, integrity, and flexibility, and lays a clear path toward a production-grade system.

## References

- [1] Elmasri, R., & Navathe, S. B. (2016). *Fundamentals of Database Systems* (7th ed.). Pearson Education.
- [2] Kent, W. (1983). *A Simple Guide to Five Normal Forms in Relational Database Theory*. Communications of the ACM, **26**(2), 120–125. <https://doi.org/10.1145/358024.358054>
- [3] Ambler, S. W. (2003). *Agile Database Techniques: Effective Strategies for the Agile Software Developer*. John Wiley & Sons.
- [4] Taylor, R. (2021). *The Evolution of Metacritic Scores: An Analysis of Aggregated Review Metrics*. Journal of Digital Media Studies, **14**(3), 145–163. <https://doi.org/10.1037/dms0000147>
- [5] GeeksforGeeks. (2024). *Normal Forms in DBMS*. <https://www.geeksforgeeks.org/normal-forms-in-dbms/>
- [6] freeCodeCamp. (2023). *Database Normalization: 1NF, 2NF, 3NF Table Examples*. <https://www.freecodecamp.org/news/database-normalization-1nf-2nf-3nf-table-examples/>
- [7] PopSQL. (2024). *Normalization in SQL DBMS: 1NF, 2NF, 3NF, and BCNF Examples*. <https://popsql.com/blog/normalization-in-sql>
- [8] StudyTonight. (2024). *Normalization in DBMS*. <https://www.studytonight.com/dbms/database-normalization.php>
- [9] Wikipedia. (2024). *Database normalization*. [https://en.wikipedia.org/wiki/Database\\_normalization](https://en.wikipedia.org/wiki/Database_normalization)